



Copyright Notice



These slides are distributed under the Creative Commons License.

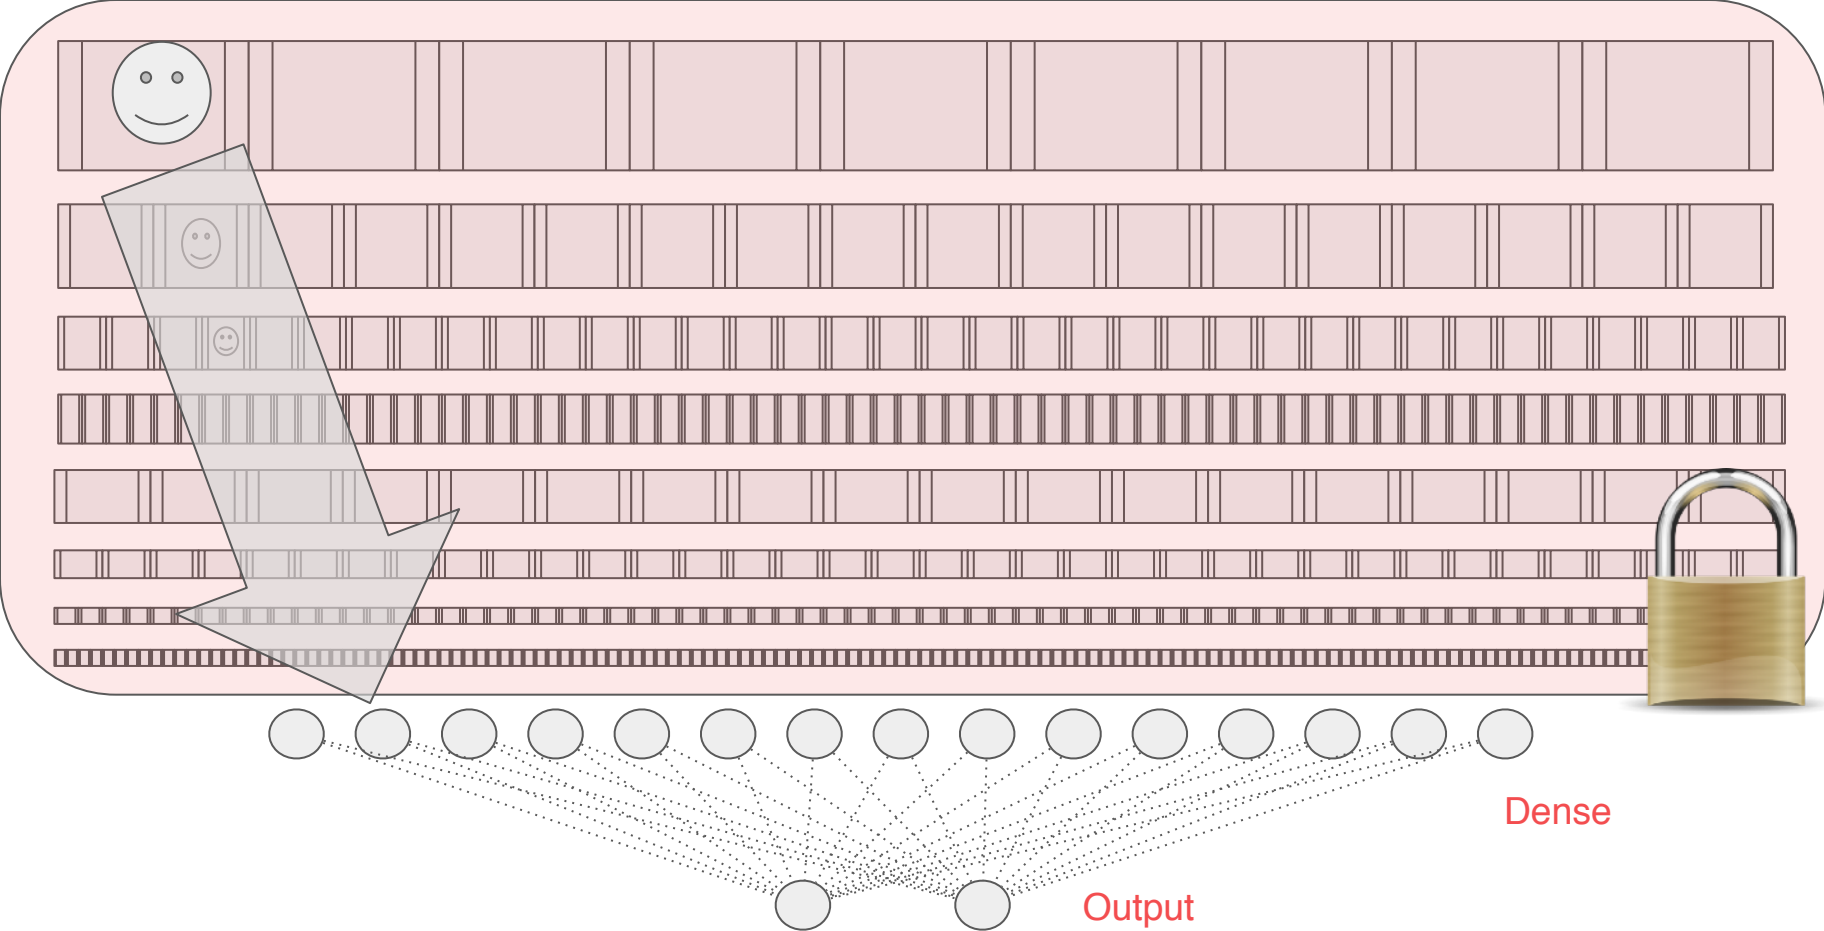
[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

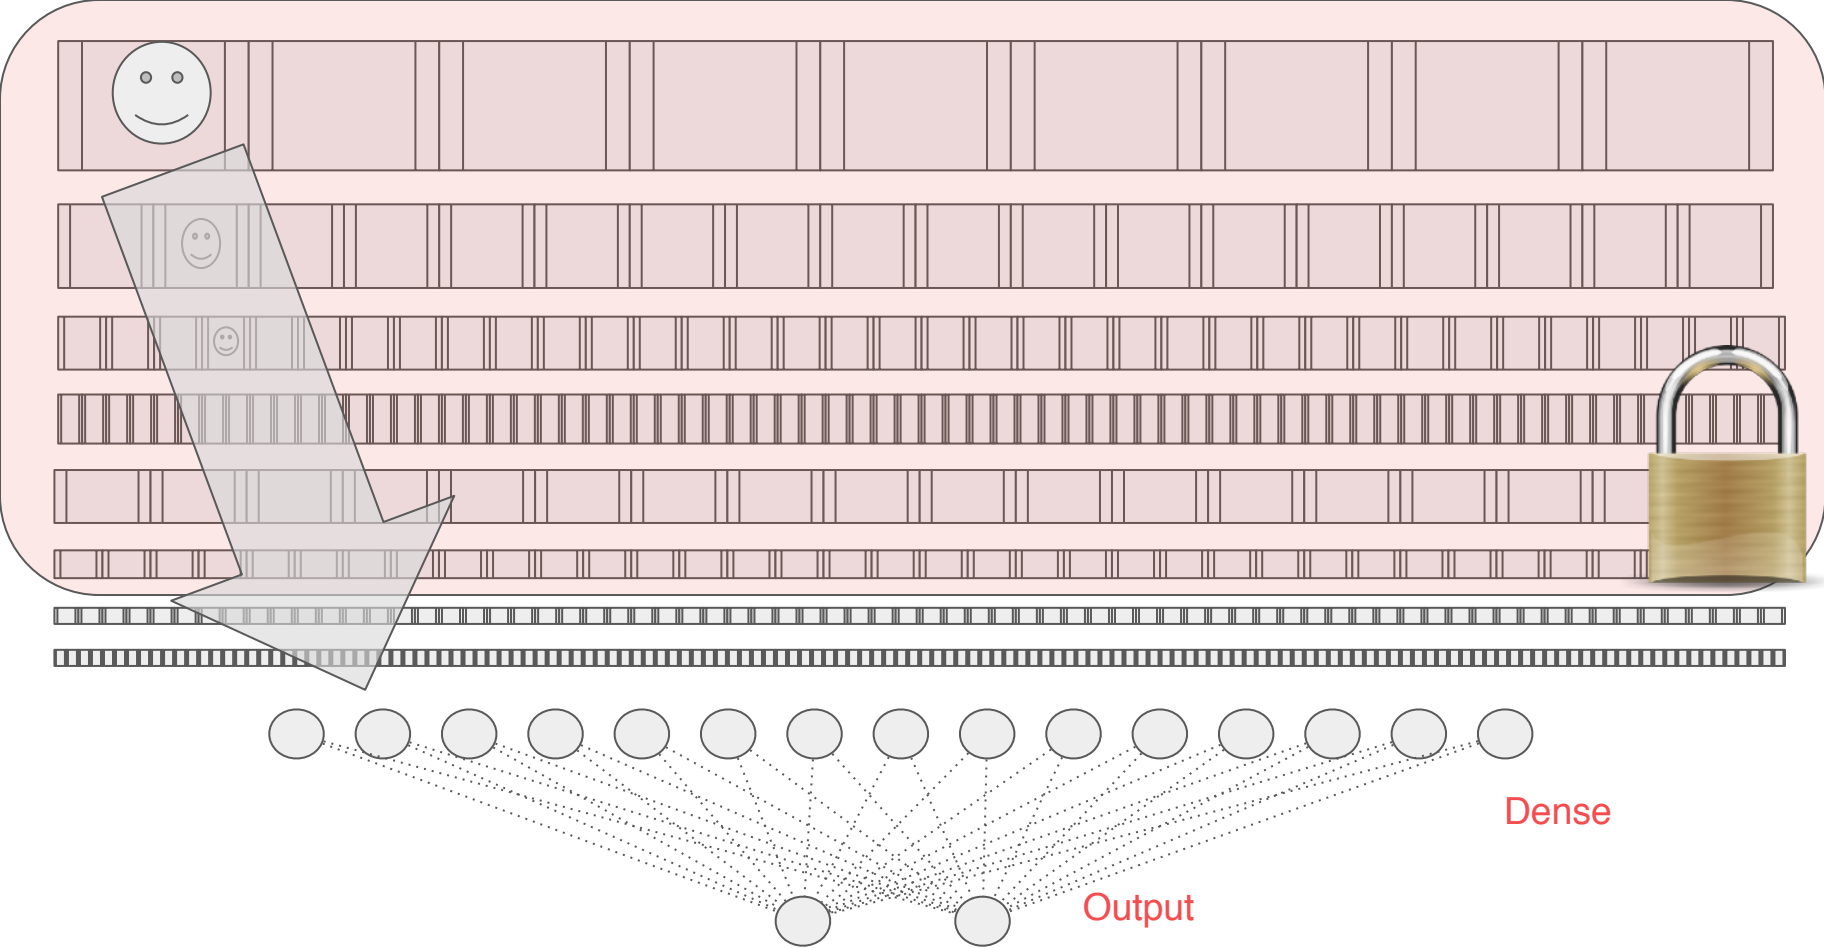
For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>













We gratefully acknowledge support from the Simons Foundation and member institutions.

arXiv.org > cs > arXiv:1512.00567

Search or Article ID

All fields



(Help | Advanced search)

Computer Science > Computer Vision and Pattern Recognition

Rethinking the Inception Architecture for Computer Vision

Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jonathon Shlens, Zbigniew Wojna

(Submitted on 2 Dec 2015 (v1), last revised 11 Dec 2015 (this version, v3))

Convolutional networks are at the core of most state-of-the-art computer vision solutions for a wide variety of tasks. Since 2014 very deep convolutional networks started to become mainstream, yielding substantial gains in various benchmarks. Although increased model size and computational cost tend to translate to immediate quality gains for most tasks (as long as enough labeled data is provided for training), computational efficiency and low parameter count are still enabling factors for various use cases such as mobile vision and big-data scenarios. Here we explore ways to scale up networks in ways that aim at utilizing the added computation as efficiently as possible by suitably factorized convolutions and aggressive regularization. We benchmark our methods on the ILSVRC 2012 classification challenge validation set demonstrate substantial gains over the state of the art: 21.2% top-1 and 5.6% top-5 error for single frame evaluation using a network with a computational cost of 5 billion multiply-adds per inference and with using less than 25 million parameters. With an ensemble of 4 models and multi-crop evaluation, we report 3.5% top-5 error on the validation set (3.6% error on the test set) and 17.3% top-1 error on the validation set.

Download:

- [PDF](#)
 - [Other formats](#)
- (license)

Current browse context:

cs.CV

[< prev](#) | [next >](#)
[new](#) | [recent](#) | [1512](#)

Change to browse by:

cs

References & Citations

- [NASA ADS](#)

<http://image-net.org/>



14,197,122 images, 21841 synsets indexed

[Explore](#) [Download](#) [Challenges](#) [Publications](#) [CoolStuff](#) [About](#)

Not logged in. [Login](#) [Signup](#)

ImageNet is an image database organized according to the **WordNet** hierarchy (currently only the nouns), in which each node of the hierarchy is depicted by hundreds and thousands of images. Currently we have an average of over five hundred images per node. We hope ImageNet will become a useful resource for researchers, educators, students and all of you who share our passion for pictures.

[Click here](#) to learn more about ImageNet, [Click here](#) to join the ImageNet mailing list.



What do these images have in common? *Find out!*

[Check out the ImageNet Challenge on Kaggle!](#)


```
import os
```

```
from tensorflow.keras import layers  
from tensorflow.keras import Model
```

[https://storage.googleapis.com/mledu-datasets/
inception_v3_weights_tf_dim_ordering_tf_kernels](https://storage.googleapis.com/mledu-datasets/inception_v3_weights_tf_dim_ordering_tf_kernels)

```
from tensorflow.keras.applications.inception_v3 import InceptionV3
```

```
local_weights_file = '/tmp/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5'
```

```
pre_trained_model = InceptionV3(input_shape = (150, 150, 3),  
                                include_top = False, => set False to go straight to Convolution  
                                weights = None)
```

```
pre_trained_model.load_weights(local_weights_file)
```

```
for layer in pre_trained_model.layers:
```

```
    layer.trainable = False => Lock all the convolutions and features that already  
                                learned into your model in all layers  
                                In order to use pre-trained model layers
```

```
pre_trained_model.summary()
```

```
last_layer = pre_trained_model.get_layer('mixed7') => Choose Layer to set the output
```

```
last_output = last_layer.output => Obtain the layer's output
```

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

=> Create Flatten Layer by taking the set output

```
x = layers.Dense(1024, activation='relu')(x)
```

```
x = layers.Dense(1, activation='sigmoid')(x)
```

```
model = Model(pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(learning_rate=0.0001),
```

```
              loss = 'binary_crossentropy',
```

```
              metrics = ['acc'])
```

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

```
x = layers.Dense(1024, activation='relu')(x)
```

=> Hidden Layers

```
x = layers.Dense(1, activation='sigmoid')(x)
```

```
model = Model(pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(learning_rate=0.0001),
```

```
              loss = 'binary_crossentropy',
```

```
              metrics = ['acc'])
```



```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

```
x = layers.Dense(1024, activation='relu')(x)
```

```
x = layers.Dense(1, activation='sigmoid')(x)
```

=> Set up Output Layer for your own model

```
model = Model(pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(learning_rate=0.0001),
```

```
              loss = 'binary_crossentropy',
```

```
              metrics = ['acc'])
```

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

```
x = layers.Dense(1024, activation='relu')(x)
```

```
x = layers.Dense (1, activation='sigmoid')(x)
```

```
model = Model( pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(learning_rate=0.0001),
```

```
    loss = 'binary_crossentropy',
```

```
    metrics = ['acc'])
```

=> Create your own model with all
set up layers and pre-train model

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

```
x = layers.Dense(1024, activation='relu')(x)
```

```
x = layers.Dense(1, activation='sigmoid')(x)
```

```
model = Model(pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(lr=0.0001),
```


```
              loss = 'binary_crossentropy',
```

```
              metrics = ['acc'])
```




```
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    batch_size = 20,  
    class_mode = 'binary',  
    target_size = (150, 150))
```

```
history = model.fit(  
    train_generator,  
    validation_data = validation_generator,  
    steps_per_epoch = 100,  
    epochs = 100,  
    validation_steps = 50,  
    verbose = 2)
```




Using dropout!



How to improve it?

However,
dropping out can break the
neural network
out of this potential bad
habit!



The idea behind it is to remove a random number of neurons in your neural network. This works very well for two reasons: The first is that neighboring neurons often end up with similar weights, which can lead to overfitting, so dropping some out at random can remove this. The second is that often a neuron can over-weigh the input from a neuron in the previous layer and can over specialize as a result.

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

```
x = layers.Dense(1024, activation='relu')(x)
```

```
x = layers.Dense (1, activation='sigmoid')(x)
```

```
model = Model( pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(lr=0.0001),
```

```
              loss = 'binary_crossentropy',
```

```
              metrics = ['acc'])
```

```
from tensorflow.keras.optimizers import RMSprop
```

```
x = layers.Flatten()(last_output)
```

```
x = layers.Dense(1024, activation='relu')(x)
```

```
x = layers.Dropout(0.2)(x)
```

 => Dropping 20% our neurons

```
x = layers.Dense(1, activation='sigmoid')(x)
```

```
model = Model(pre_trained_model.input, x)
```

```
model.compile(optimizer = RMSprop(lr=0.0001),
```

```
loss = 'binary_crossentropy',
```

```
metrics = ['acc'])
```

